

Exam 1 Review Problems — Solutions

STAT 251 SECTION 01

2026-09-16

Worked solutions. Answers are shown in red.

These problems cover everything on Exam 1: Weeks 1 through 3, up to and including cumulative distributions. They are **more computational than the exam** by design. The exam asks you to interpret values, identify equations and read plots; these ask you to work the arithmetic behind them.

Work them by hand. Every dataset is small enough to do on paper, and a calculator is fine.

Problem 1 — The anatomy of data

A researcher recorded the following information for each of six trees in a city park.

Tree ID	Species	Health rating	Branches removed	Trunk width (cm)
T-01	Maple	Good	3	31.4
T-02	Oak	Fair	0	58.2
T-03	Maple	Excellent	5	22.7
T-04	Pine	Good	1	44.0
T-05	Oak	Poor	8	61.9
T-06	Maple	Good	2	27.5

(a) How many observations are in this data table, and how many variables?

6 observations and 5 variables

(b) What is the value of n for this sample?

$n = 6$

(c) Classify each of the four measured variables as nominal, ordinal, discrete or continuous.

Variable	Your answer
Species	nominal
Health rating	ordinal
Branches removed	discrete
Trunk width (cm)	continuous

Species is a category with no natural ordering, so nominal. **Health rating** is a category *with* an ordering – poor, fair, good, excellent – so ordinal. **Branches removed** counts things, so it takes whole-number values only: discrete. **Trunk width** is a measurement on a scale, recorded to a tenth of a centimetre, so continuous.

Note that *Tree ID* is a label, not a measured variable – it identifies the observation rather than describing it.

(d) These six trees are a sample. Describe the population the researcher is presumably interested in, and say whether N is known.

The population is all of the trees in the city park (or, more broadly, all trees the park service manages). N is the number of trees in that population. Here it is not given – and unlike most populations we will meet, this one is finite and could actually be counted.

Problem 2 — Statistics, parameters, and inference

A city has 8,500 registered bicycles. A researcher takes a random sample of 200 of them and finds that 46 have no working lights. She also measures the age of each sampled bicycle and finds a mean age of 6.2 years.

(a) Give the values of N and n .

$$N = 8,500 \text{ and } n = 200$$

(b) Compute the sample proportion of bicycles with no working lights. Which symbol denotes it?

$$\hat{p} = 46/200 = 0.23$$

The **sample** proportion is written $\hat{p}$. The hat marks it as a quantity computed from a sample. The corresponding population proportion – the fraction of all 8,500 bicycles with no working lights – is written p and is unknown.

(c) For each quantity, say whether it is a **statistic** or a **parameter**, and give the symbol that denotes it.

Quantity	Statistic or parameter	Symbol
The mean age of the 200 sampled bicycles	statistic	$\bar{x}$
The mean age of all 8,500 registered bicycles	parameter	μ
The proportion of the 200 sampled with no lights	statistic	$\hat{p}$
The proportion of all 8,500 with no lights	parameter	p

(d) The researcher reports the mean age of the 200 bicycles she measured, then uses it to estimate the mean age of all 8,500. Which of those two acts is descriptive statistics and which is inferential? Explain what makes the difference.

Reporting the mean of the 200 measured bicycles is **descriptive**: it summarizes the data in hand and claims nothing beyond it. Using that mean to estimate the mean age of all 8,500 is **inferential**: it reaches past the sample to a population that was never fully observed.

The difference is not which data you have – both acts use the same 200 bicycles. It is *what you are claiming*: whether the conclusion stops at the sample or extends beyond it.

Problem 3 — Reading the formulas

On the exam the formulas are provided **unlabeled** — the equations only, with no names. This problem is practice at telling them apart.

(a) Each equation below computes one of the quantities in the list. Write the name of the quantity next to its equation.

sample mean · population mean · sample proportion · sample variance · sample standard deviation · interquartile range

Equation	What it computes	Symbol
$\frac{1}{n} \sum_{i=1}^n x_i$	sample mean	$\bar{x}$
$\sqrt{\frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2}$	sample standard deviation	s
$Q3 - Q1$	interquartile range	IQR
$\frac{1}{N} \sum_{i=1}^N x_i$	population mean	μ
$\frac{\text{number of observations in the category}}{n}$	sample proportion	$\hat{p}$
$\frac{1}{n-1} \sum_{i=1}^n (x_i - \bar{x})^2$	sample variance	s^2

Three things give each one away. The **divisor**: n for a sample mean, N for a population mean, $n - 1$ for a variance. The **square root**, which is the only difference between s^2 and s . And whether the deviations are **squared** at all.

The two that get confused most often are the middle pair: s^2 and s are the same expression apart from the root, so check for it every time.

(b) (notation) Answer each in a few words.

Question	Your answer
What does n stand for?	the number of observations in the sample
What does N stand for?	the number of observations in the population
What does the <i>hat</i> on $\hat{p}$ tell you?	that it is computed from a sample, and estimates the parameter p
In $\sum_{i=1}^n x_i$, what does the $i = 1$ below the Σ mean?	start the count at the first observation and add up to the n^{th}

(c) Consider the following sample of five observations of a quantitative variable X :

$$X = \{2, 5, 6, 9, 13\}$$

Show how to compute the mean $\bar{x}$ and the standard deviation s .

Mean. $\bar{x} = \frac{1}{n} \sum_{i=1}^n x_i = \frac{2+5+6+9+13}{5} = \frac{35}{5} = 7$

Standard deviation. Take the deviation of each observation from the mean, square it, and add:

$$\sum_{i=1}^n (x_i - \bar{x})^2 = (-5)^2 + (-2)^2 + (-1)^2 + 2^2 + 6^2 = 25 + 4 + 1 + 4 + 36 = 70$$

Divide by $n - 1$ to get the variance, then take the square root:

$$s^2 = \frac{70}{5-1} = \frac{70}{4} = 17.5 \quad s = \sqrt{17.5} \approx 4.18$$

Note the deviations themselves – $-5, -2, -1, 2, 6$ – add to 0. That happens for **every** dataset, because the mean is the balancing point, and it is why the deviations must be squared before they are added.

(d) Now consider the following sample of ten observations of a qualitative variable Y , the way each of ten students travelled to campus:

$$Y = \{\text{Bus, Walk, Bike, Bus, Walk, Bus, Car, Bus, Walk, Bike}\}$$

Show how to compute the proportion of observations in the category **Bus**.

Count the observations in the category, then divide by the sample size:

$$\hat{p} = \frac{\text{number of observations in the category}}{n} = \frac{4}{10} = 0.4$$

So 40% of the sample travelled by bus. Note that a proportion is the one summary here that *can* be computed for a qualitative variable – a mean and a standard deviation cannot, because the categories are not numbers.

Problem 4 — The distribution of a qualitative variable

A campus coffee stand recorded the drink ordered by each of 20 customers one morning:

$$X = \{\text{Latte, Drip, Mocha, Latte, Drip, Espresso, Latte, Drip, Latte, Mocha, Drip, Latte, Espresso, Drip, Mocha, Latte, Drip, Latte, Mocha, Latte}\}$$


(a) Complete the frequency and relative frequency table.

Drink	Frequency $F(x)$	Relative frequency $RF(x)$
Drip	6	0.30
Espresso	2	0.10
Latte	8	0.40
Mocha	4	0.20

(b) What is the modal category?

Latte, with a frequency of 8

(c) Draw a bar graph of the distribution on the axes below.



(d) Would it make sense to add a *cumulative relative frequency* column to the table in part (a)? Explain.

No. Cumulative relative frequency is the proportion of observations *less than or equal to a given value*, and that phrase only means something if the values can be put in order. Drink type is a **nominal** variable – there is no sense in which a mocha is ‘less than’ a latte – so there is no meaningful order to accumulate along.

Had the variable been ordinal, say drink size (small, medium, large), a cumulative relative frequency column would be perfectly sensible.

Problem 5 — Frequency tables and the mean from a table

Twenty-five households were asked how many vehicles they own. The responses are summarized below.

Vehicles x	Frequency $F(x)$	Relative frequency $RF(x)$	Cumulative relative frequency $CRF(x)$
0	3	0.12	0.12
1	9	0.36	0.48
2	8	0.32	0.80
3	4	0.16	0.96
4	1	0.04	1.00

(a) Complete the relative frequency and cumulative relative frequency columns.

(b) Compute the mean number of vehicles using the **frequency** formula, $\bar{x} = \frac{1}{n} \sum_x x F(x)$.

$$\bar{x} = \frac{0(3)+1(9)+2(8)+3(4)+4(1)}{25} = \frac{0+9+16+12+4}{25} = \frac{41}{25} = 1.64 \text{ vehicles.}$$

(c) Now compute it again using the **relative frequency** formula, $\bar{x} = \sum_x x RF(x)$, and confirm you get the same answer.

$$\bar{x} = 0(0.12) + 1(0.36) + 2(0.32) + 3(0.16) + 4(0.04) = 0 + 0.36 + 0.64 + 0.48 + 0.16 = 1.64 \text{ vehicles.}$$

The two formulas are the same calculation. Dividing by n at the end (the frequency version) and dividing by n at the start to build the relative frequencies (this version) get you to the same place. This is also why you must not divide by n again in the relative frequency version.

(d) What proportion of these households own 2 or fewer vehicles? Which column did you read that from?

0.80, read from the cumulative relative frequency column at $x = 2$

(e) What is the median number of vehicles owned? (You can read this off the cumulative relative frequency column.)

The median is the 50th percentile. Reading down the CRF column, it passes 0.50 at $x = 2$ (the CRF is 0.48 at $x = 1$ and 0.80 at $x = 2$), so the median is 2 vehicles.

Check it against the raw data if you like: with $n = 25$ the median is the 13^{th} value in order, and values 13 through 20 are all 2s.

Problem 6 — Binning a continuous variable

The following are the times, in seconds, that 16 users took to complete a task on a website.

$$X = \{12.4, 14.1, 15.8, 16.2, 17.5, 18.0, 18.3, 19.1, 20.2, 20.6, 21.7, 22.4, 23.9, 25.1, 26.8, 28.4\}$$

(a) Using the **square root rule**, how many bins k should you use for a histogram of these 16 observations?

$$k = \text{round}(\sqrt{16}) = 4 \text{ bins}$$


(b) Compute the bin width w implied by that choice, using $w = \frac{\max x - \min x}{k}$.

$$w = (28.4 - 12.4)/4 = 16/4 = 4.0 \text{ seconds}$$

(c) Complete the frequency table using the intervals given.

Time (seconds)	Frequency $F(x)$	Relative frequency $RF(x)$
$12 < x \leq 16$	3	0.1875
$16 < x \leq 20$	5	0.3125
$20 < x \leq 24$	5	0.3125
$24 < x \leq 28.5$	3	0.1875

(d) Draw the histogram on the axes below.



(e) Describe the shape of this distribution.

Roughly **symmetric and unimodal**, with the bulk of the observations in the two middle bins and the counts tapering off evenly at both ends. There is no long tail in either direction and no obvious gap, so nothing here looks like an outlier.

(f) Give one thing you can determine from the table in part (c), and one thing you **cannot**, that you could have determined from the raw data.

Can: the proportion of users finishing in 20 seconds or less ($8/16 = 0.5$), or the number in any one interval. **Cannot:** the fastest or slowest individual time, the exact mean, or the median – once the values are grouped, the individual measurements are gone.

That loss of precision is the cost of binning. It buys a readable summary of the shape, which the raw list of 16 numbers does not give you.

Problem 7 — Dot plots and stem-and-leaf plots


A biologist recorded the number of eggs in each of 15 birds' nests:

$$X = \{21, 24, 24, 27, 28, 31, 33, 33, 33, 35, 38, 41, 42, 44, 47\}$$

(a) Construct a **stem-and-leaf plot**, using the tens digit as the stem.

2		1	4	4	7	8	
3		1	3	3	3	5	8
4		1	2	4	7		

(b) Construct a **dot plot** on the number line below.



(c) What is the mode of this variable?

33 eggs, which occurs three times

(d) Is the mode a measure of the center of this distribution? Explain.

No – the mode is a measure of **location**. It tells you where the peak is, i.e. the most frequent value, not where the center is. Here the distribution happens to be roughly symmetric so the mode of 33 does land near the middle, but that is a coincidence of this data rather than a property of the mode.

For a skewed distribution the mode can sit well away from the center, and a bimodal distribution has two modes with the center usually falling in the gap between them, where the data are sparsest.

Problem 8 — Center, and what an outlier does to it

Consider the following eight observations of a variable X , the number of minutes eight customers spent in a shop:

$$X = \{4, 6, 6, 9, 11, 12, 14, 16\}$$

(a) Compute the mean, the median and the mode.

Statistic	Your answer
Mean $\bar{x}$	$78/8 = 9.75$ minutes
Median	$(9 + 11)/2 = 10$ minutes
Mode	6 minutes

(b) A ninth customer stays for **78 minutes**. Recompute the mean and the median with that observation included.

$$\bar{x} = (78 + 78)/9 = 156/9 \approx 17.33 \text{ minutes.}$$

The median is now the 5th of 9 ordered values, which is **11 minutes**.

(c) Which statistic moved more, and why?

The **mean** moved far more – from 9.75 up to about 17.33, an increase of roughly 7.6 minutes – while the median moved from 10 to 11, just one minute.

The mean uses the *value* of every observation, so a single very large value pulls it upward in proportion to how extreme it is. The median depends only on *position* in the ordered data; adding one more value at the top shifts the middle position over by one place and nothing more.

Notice too that after the outlier the mean of 17.33 is larger than *every observation except the outlier itself*. It is no longer describing a typical customer at all.

(d) Which of the two is **resistant**? Give a one-sentence definition of that word.

The median is resistant.

A summary measure is **resistant** when a single extreme observation barely moves it. Whether a measure is resistant comes down to one thing: whether it uses the values of every observation, or only positions in the ordered data.

Problem 9 — Percentiles and quartiles

The following are the ages of the 11 members of a community board, sorted:

$$X = \{29, 34, 38, 41, 45, 49, 52, 56, 61, 67, 73\}$$

(a) Find the median, Q_2 .

$Q2 = 49$, the 6th of 11 ordered values

(b) Find $Q1$ and $Q3$ using the **method of halves** we use in this class. Remember that when n is odd the median is **left out** of both halves.

Dropping the median, the lower half is {29, 34, 38, 41, 45} and the upper half is {52, 56, 61, 67, 73}. Each has five values, so each median is its middle value:

$Q1 = 38$ years and $Q3 = 61$ years.

(c) A twelfth member, aged 58, joins the board. Find $Q1$, $Q2$ and $Q3$ for the new set of 12 ages.

The ordered ages are now {29, 34, 38, 41, 45, 49, 52, 56, 58, 61, 67, 73}.

With $n = 12$ there is no single middle value, so the median is the average of the 6th and 7th ages:
 $Q2 = (49 + 52)/2 = 50.5$ years. Because n is even there is nothing to leave out – the data splits cleanly into two halves of six.

The lower half is {29, 34, 38, 41, 45, 49}, so $Q1 = (38 + 41)/2 = 39.5$ years.

The upper half is {52, 56, 58, 61, 67, 73}, so $Q3 = (58 + 61)/2 = 59.5$ years.

Compare this with part (b): the method is the same both times. The only thing that changes is that an odd n leaves you a middle value to drop, and an even n does not.

(d) What is the IQR , and what does it measure?

$IQR = 61 - 38 = 23$ years

The IQR is the distance between $Q1$ and $Q3$, so it describes how spread out the **middle 50%** of the data are. Because it ignores the extreme ends of the distribution entirely, it is highly resistant to outliers.

(e) The board member aged 67 is at roughly what percentile? Answer in a sentence that says what that means.

Ten of the 11 members are aged 67 or younger, so 67 sits at about the $10/11 \approx 91^{st}$ percentile: roughly 91% of the board is at or below that age, and only one member is older.

Problem 10 — Deviations, variance and standard deviation

A quality inspector measured the weight, in grams, of six packets from a filling machine:

$$X = \{48, 51, 52, 55, 57, 59\}$$

(a) Compute the sample mean $\bar{x}$.

$$\bar{x} = 322/6 \approx 53.67 \text{ grams}$$

(b) Complete the deviation table below, then add up the deviation column.

x_i	$x_i - \bar{x}$	$(x_i - \bar{x})^2$
48	-5.67	32.11
51	-2.67	7.11
52	-1.67	2.78
55	1.33	1.78
57	3.33	11.11
59	5.33	28.44

The deviation column sums to **0** (up to rounding), and the squared deviation column sums to about **83.33**.

The deviations summing to zero is not a property of these six numbers – it happens for every dataset, because the mean is the balancing point. That is exactly why the variance squares the deviations before adding them: without squaring, the positives and negatives always cancel and you learn nothing.

(c) Compute the sample variance s^2 and the sample standard deviation s .

$$s^2 = \frac{83.33}{6-1} = \frac{83.33}{5} \approx 16.67$$

$$s = \sqrt{16.67} \approx 4.08 \text{ grams}$$

(d) State the units of $\bar{x}$, of s^2 and of s .

Quantity	Units
$\bar{x}$	grams
s^2	grams squared
s	grams

Squaring the deviations squares their units too, which is why the variance is in grams squared – a quantity nobody can picture. Taking the square root undoes it and puts s back on the scale of the original data, where it reads as roughly the typical distance of a packet from the mean weight.

(e) Compute the range. Which is the better description of the spread of this data, the range or the standard deviation, and why?

$$\text{Range} = 59 - 48 = 11 \text{ grams.}$$

The **standard deviation** is the better description here. The range is built from the two most extreme observations and *nothing else*, so it throws away the other four measurements entirely and moves by the full amount if either endpoint changes. The standard deviation uses every observation.

The range is worth reporting as a quick descriptive fact, but it is almost never used as *the* measure of spread.

Problem 11 — Five-number summary, boxplots and outliers

A clinic recorded the number of days between a referral and an appointment for 12 patients:

$$X = \{3, 5, 6, 8, 9, 11, 12, 14, 15, 17, 19, 58\}$$

(a) Give the five-number summary. Use the method of halves for the quartiles.

Minimum	Q1	Median	Q3	Maximum
3	7	11.5	16	58

$n = 12$ is even, so there is no middle value to leave out: the median is the average of the 6th and 7th values, $Q2 = (11 + 12)/2 = 11.5$ days, and the data splits cleanly into two halves of six.

The lower half is $\{3, 5, 6, 8, 9, 11\}$ with median $Q1 = (6 + 8)/2 = 7$; the upper half is $\{12, 14, 15, 17, 19, 58\}$ with median $Q3 = (15 + 17)/2 = 16$.

(b) Compute the IQR and both outlier boundaries.

$$IQR = 16 - 7 = 9 \text{ days, so } 1.5 \times IQR = 13.5.$$

$$\text{Lower boundary} = 7 - 13.5 = -6.5 \text{ days.}$$


Upper boundary = $16 + 13.5 = 29.5$ days.

Nothing can fall below -6.5 days, so only the upper boundary can flag anything here.

(c) Which observations, if any, does the $1.5 \times IQR$ rule flag as outliers?

Just one: the 58-day wait, which is above the upper boundary of 29.5

(d) Draw the boxplot on the axis below. Be careful about where the right whisker ends.



The right whisker stops at **19**, the largest observation that is still at or below the upper boundary of 29.5 – *not* at the maximum of 58, and not at the boundary itself. A whisker always ends on an actual data value. The 58 is drawn as its own point beyond the whisker.

The left whisker runs all the way out to the minimum of 3, because nothing is flagged at the low end.

(e) Describe the shape of this distribution from the boxplot, and say what that implies about the mean relative to the median.

The distribution is **skewed right**: the median sits left of centre in the box, the right whisker is longer than the left, and there is a high outlier.

A right skew pulls the mean toward the long tail, so the **mean will be greater than the median**. The mean here is 14.75 days against a median of 11.5.

(f) The clinic manager wants to delete the 58-day wait because it is “clearly a mistake.” What would you tell her?

That being flagged by the $1.5 \times IQR$ rule is an **arithmetic result, not a verdict**. The rule reports only that a value sits unusually far from the middle half of the data; it cannot distinguish a recording error from a genuinely unusual case, and it marks the observation as worth a second look rather than

as wrong.

A 58-day wait is entirely plausible – it may be the single most important number in the dataset for a manager trying to improve scheduling. Deleting it because a rule flagged it would throw away real information.


Problem 12 — Cumulative distributions

A librarian recorded the number of books checked out by each of 20 patrons in one hour.

(a) Complete the table.

Books x	Frequency $F(x)$	Relative frequency $RF(x)$	Cumulative relative frequency $CRF(x)$
0	2	0.10	0.10
1	5	0.25	0.35
2	6	0.30	0.65
3	4	0.20	0.85
4	2	0.10	0.95
5	1	0.05	1.00

(b) Plot the cumulative distribution on the axes below. Remember it is a **step function**: flat between observed values, jumping at each one.



(c) Use the plot to find $Q1$, the median, and $Q3$.

Read across from 0.25, 0.50 and 0.75 on the y -axis to the curve, then drop to the x -axis. Each of the three readings crosses a jump:

$Q1 = 1$ book, $Q2 = 2$ books, $Q3 = 3$ books.

Check them against halving the ordered data – they agree. The cumulative distribution is not a new definition of a quartile, just a faster way to read one off, and unlike halving it gives you *any* percentile.

(d) What proportion of patrons checked out 2 or fewer books?

0.65, read straight off the CRF column at $x = 2$

(e) Describe the shape of this distribution, and say whether the mean will be greater than, less than, or equal to the median.

The distribution is **skewed right**: the frequencies rise to a peak at 2 books and then trail off toward the larger values.

A right skew pulls the mean toward the long tail while leaving the median where it is, so the **mean is greater than the median**. Here the mean is $42/20 = 2.1$ books against a median of 2.

A note on the exam

These solutions do more arithmetic than the exam requires. Exam 1 asks you to interpret values, identify equations and read plots. The formulas are shown during the exam, but **unlabeled** – identifying which is which is on you.